

Search  
spaces  
12 sept  
2022

CSCI  
373

# train

rules

- change one letter at a time to get from train to prawn
- each intermediate word must be valid english

# prawn



# train

rules

- change one letter at a time to get from train to prawn
- each intermediate word must be valid english

# prawn

your score  
is the number  
of intermediate  
words -



train

trail

trawl

brawl

brawn

prawn



train

your answer  
here

prawn

Can you  
do better?



train

brain

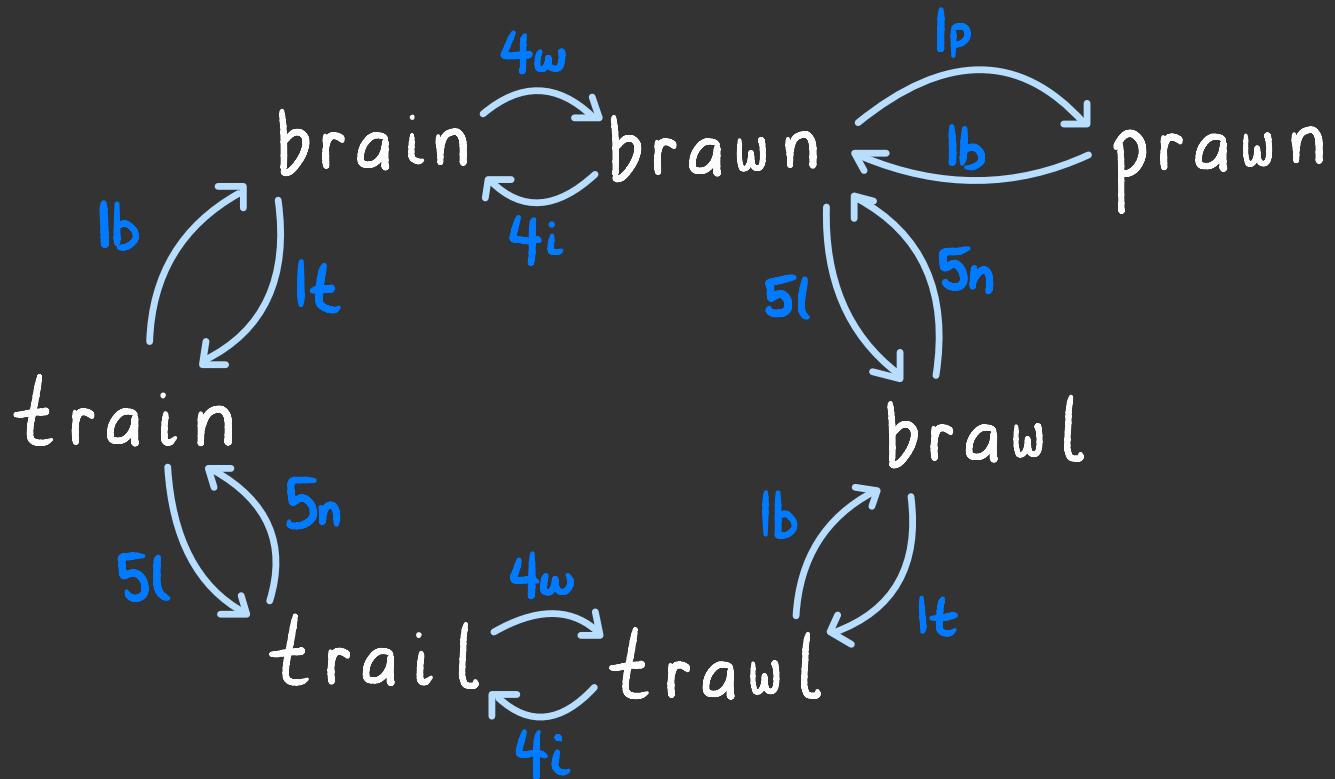
brown

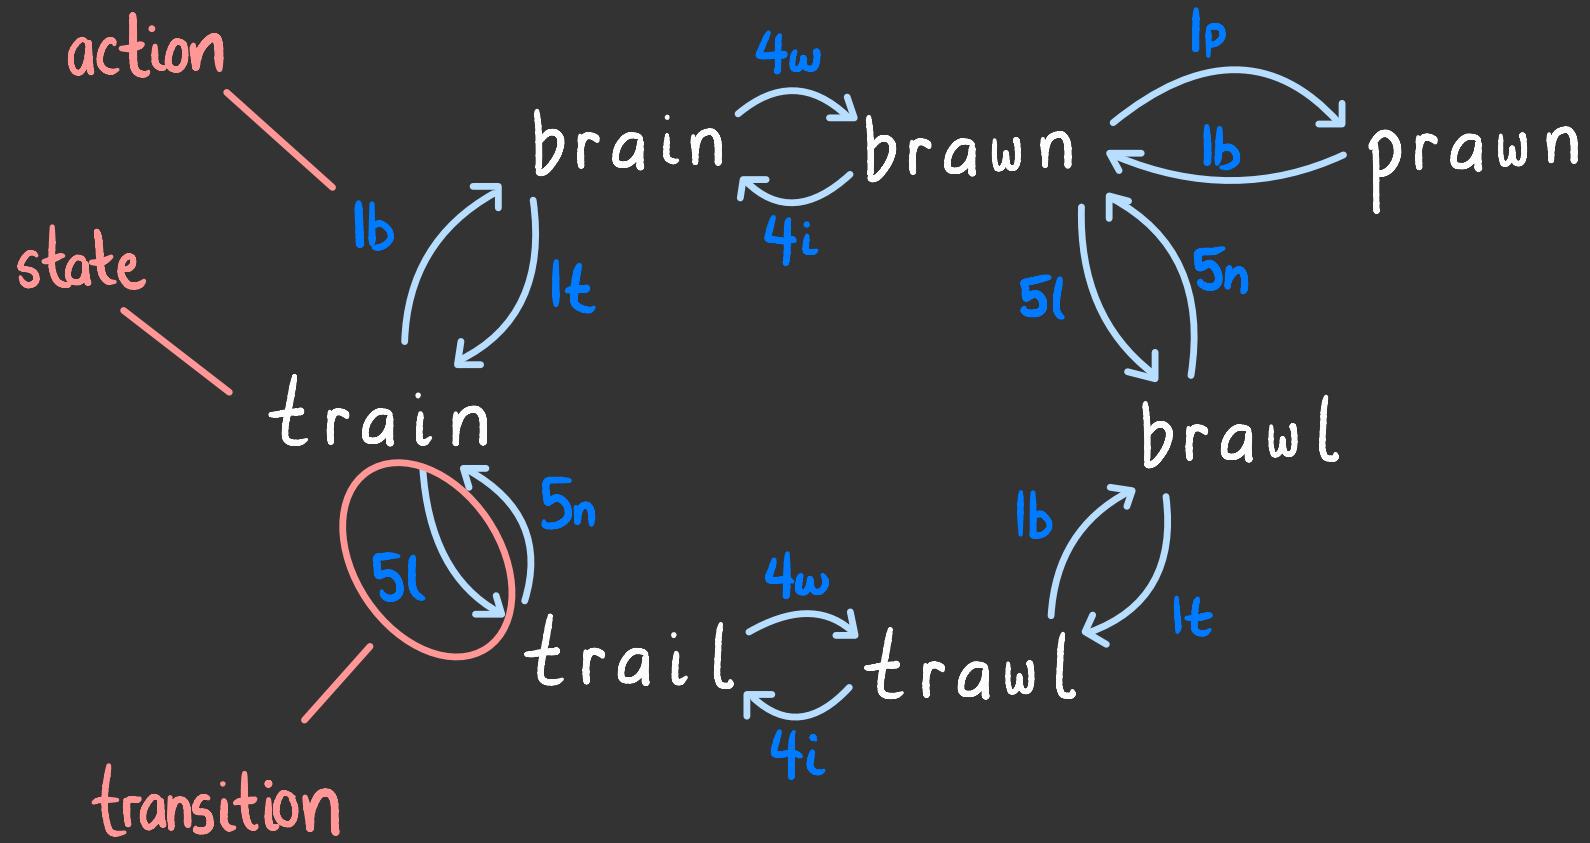
prawn

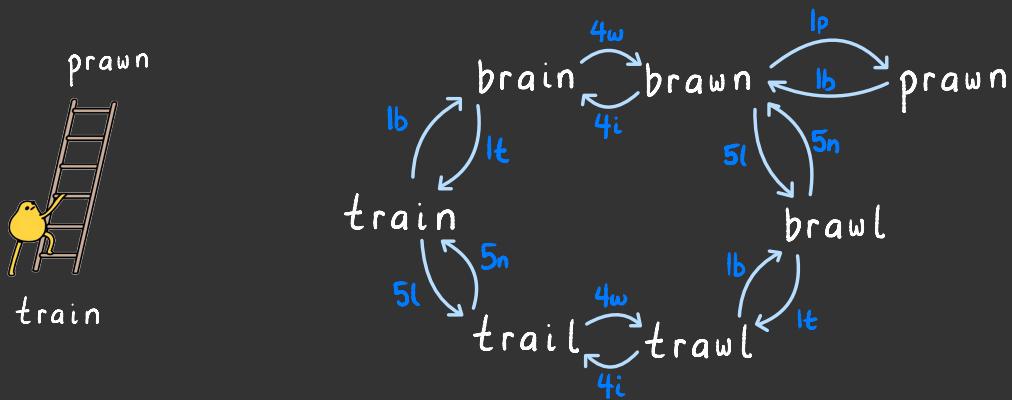
Can you  
do better?



how do we describe  
word ladder?

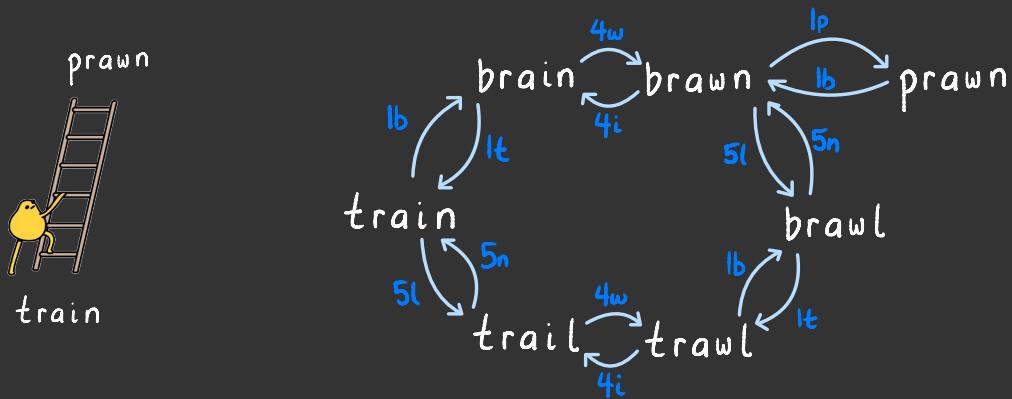






define a state machine as a tuple  $(Q, \Sigma, \Delta, q_0, F)$  where:

- $Q$  is a set of states
- $\Sigma$  is a set of actions
- $\Delta \subseteq Q \times \Sigma \times Q$  is a set of permitted transitions
- $q_0 \in Q$  is the initial state
- $F \subseteq Q$  is a set of final states

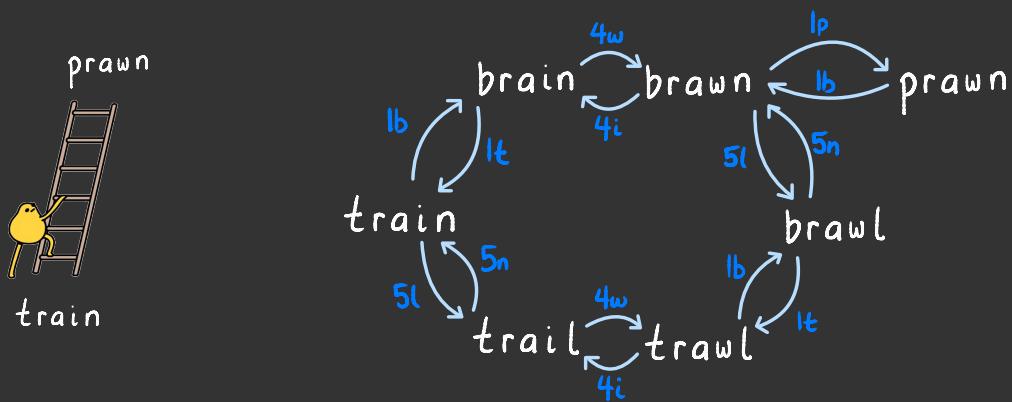


define a state machine as a tuple  $(Q, \Sigma, \Delta, q_0, F)$  where:

- $Q$  is a set of states
- $\Sigma$  is a set of actions
- $\Delta \subseteq Q \times \Sigma \times Q$  is a set of permitted transitions
- $q_0 \in Q$  is the initial state
- $F \subseteq Q$  is a set of final states

A search problem can be defined formally as follows:

- A set of possible states that the environment can be in. We call this the **state space**.
- The **initial state** that the agent starts in. For example: *Arad*.
- A set of one or more **goal states**. Sometimes there is one goal state (e.g., *Bucharest*), sometimes there is a small set of alternative goal states, and sometimes the goal is defined by a property that applies to many states (potentially an infinite number). For example, in a vacuum-cleaner world, the goal might be to have no dirt in any location, regardless of any other facts about the state. We can account for all three of these possibilities by specifying an Is-Goal method for a problem. In this chapter we will sometimes say “the goal” for simplicity, but what we say also applies to “any one of the possible goal states.”
- The **actions** available to the agent. Given a state  $s$ ,  $\text{ACTIONS}(s)$  returns a finite<sup>2</sup> set of actions that can be executed in  $s$ . We say that each of these actions is **applicable** in  $s$ . An example:  
 $\text{ACTIONS}(\text{Arad}) = \{\text{ToSibiu}, \text{ToTimisoara}, \text{ToZerind}\}$ .
- A **transition model**, which describes what each action does.  $\text{RESULT}(s, a)$  returns the state that results from doing action  $a$  in state  $s$ . For example,  
 $\text{RESULT}(\text{Arad}, \text{ToZerind}) = \text{Zerind}$ .
- An **action cost function**, denoted by  $\text{ACTION-COST}(s, a, s')$  when we are programming or  $c(s, a, s')$  when we are doing math, that gives the numeric cost of applying action  $a$  in state  $s$  to reach state  $s'$ . A problem-solving agent should use a cost function that reflects its own performance measure; for example, for route-finding agents, the cost of an action might be the length in miles (as seen in Figure 3.1), or it might be the time it takes to complete the action.



define a state machine as a tuple  $(Q, \Sigma, \Delta, q_0, F)$  where:

- $Q$  is a set of states
- $\Sigma$  is a set of actions
- $\Delta \subseteq Q \times \Sigma \times Q$  is a set of permitted transitions
- $q_0 \in Q$  is the initial state
- $F \subseteq Q$  is a set of final states

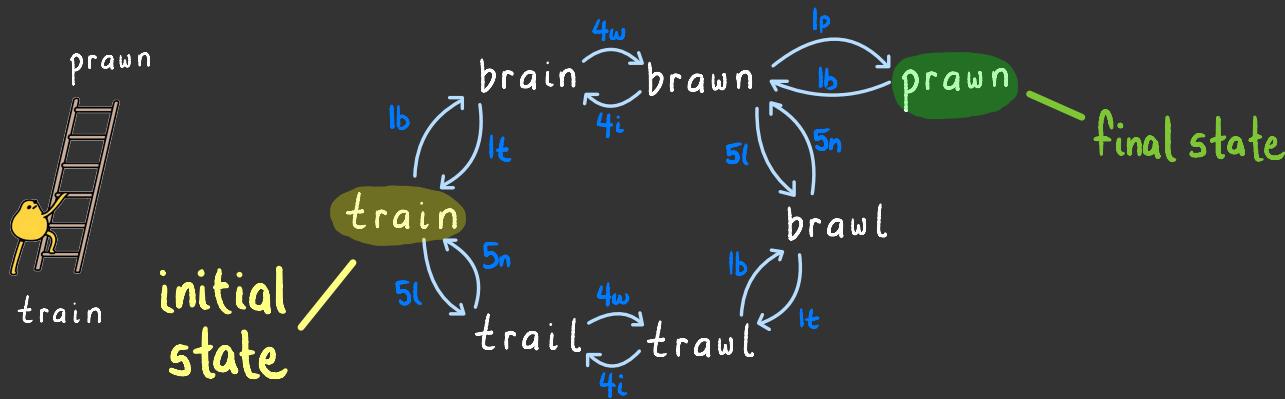
$$Q = \{ \text{train}, \text{brain}, \text{trail}, \dots \}$$

$$\Sigma = \{ 1a, \dots, 5z \}$$

$$\Delta = \{ (\text{train}, 1b, \text{brain}), \dots \}$$

$$q_0 = \text{train}$$

$$F = \{ \text{prawn} \}$$



define a state machine as a tuple  $(Q, \Sigma, \Delta, q_0, F)$  where:

- $Q$  is a set of states
- $\Sigma$  is a set of actions
- $\Delta \subseteq Q \times \Sigma \times Q$  is a set of permitted transitions
- $q_0 \in Q$  is the initial state
- $F \subseteq Q$  is a set of final states

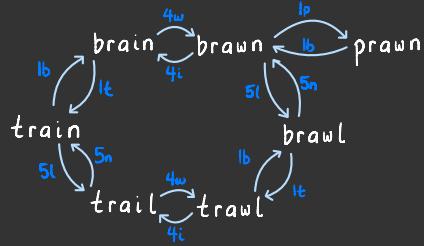
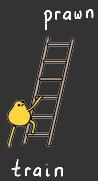
$$Q = \{ \text{train}, \text{brain}, \text{trail}, \dots \}$$

$$\Sigma = \{ 1a, \dots, 5z \}$$

$$\Delta = \{ (\text{train}, 1b, \text{brain}), \dots \}$$

$$q_0 = \text{train}$$

$$F = \{ \text{prawn} \}$$



define a **state machine** as a tuple  $(Q, \Sigma, \Delta, q_0, F)$  where:

- $Q$  is a set of **states**
- $\Sigma$  is a set of **actions**
- $\Delta \subseteq Q \times \Sigma \times Q$  is a set of permitted **transitions**
- $q_0 \in Q$  is the **initial state**
- $F \subseteq Q$  is a set of **final states**

$$Q = \{\text{train}, \text{brain}, \text{trail}, \dots\}$$

$$\Sigma = \{1a, \dots, 5z\}$$

$$\Delta = \{(\text{train}, 1b, \text{brain}), \dots\}$$

$$q_0 = \text{train}$$

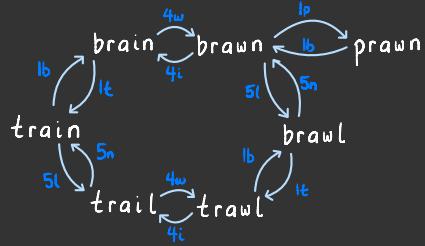
$$F = \{\text{prawn}\}$$

define a **weighted state machine** as a tuple  $(Q, \Sigma, \Delta, q_0, F, w)$  where:

- $(Q, \Sigma, \Delta, q_0, F)$  is a state machine
- $w: \Delta \rightarrow \mathbb{R}$  assigns a real-valued weight to each transition

what is the weight function  $w$  for word ladder?

your answer  
here



define a **state machine** as a tuple  $(Q, \Sigma, \Delta, q_0, F)$  where:

- $Q$  is a set of **states**
- $\Sigma$  is a set of **actions**
- $\Delta \subseteq Q \times \Sigma \times Q$  is a set of permitted **transitions**
- $q_0 \in Q$  is the **initial state**
- $F \subseteq Q$  is a set of **final states**

$$Q = \{\text{train}, \text{brain}, \text{trail}, \dots\}$$

$$\Sigma = \{1a, \dots, 5z\}$$

$$\Delta = \{(\text{train}, 1b, \text{brain}), \dots\}$$

$$q_0 = \text{train}$$

$$F = \{\text{prawn}\}$$

define a **weighted state machine** as a tuple  $(Q, \Sigma, \Delta, q_0, F, w)$  where:

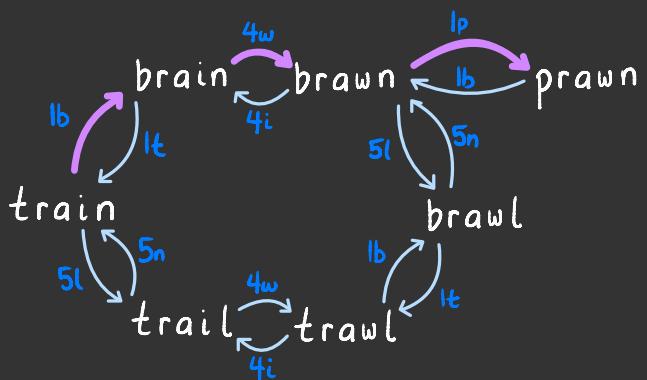
- $(Q, \Sigma, \Delta, q_0, F)$  is a state machine
- $w: \Delta \rightarrow \mathbb{R}$  assigns a real-valued weight to each transition

what is the weight function  $w$  for word **ladder**?

$$w = \{\delta \mapsto 1 \mid \delta \in \Delta\}$$

# search paths

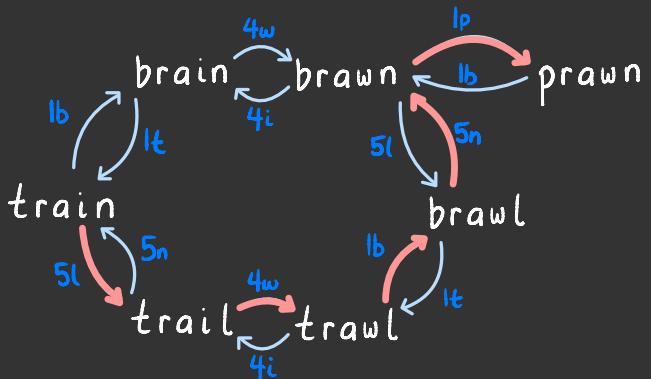
train  
brain  
brawn  
prawn



formally:

your answer here

train  
trail  
trawl  
brawl  
brawn  
prawn

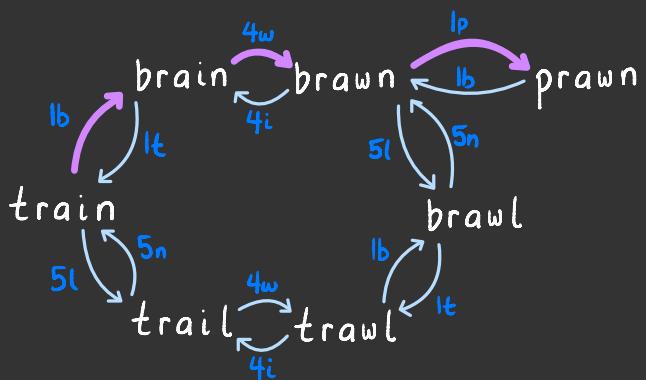


formally:

your answer here

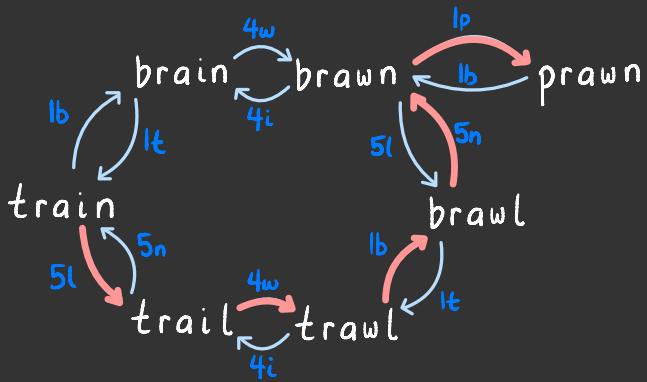
# search paths

train  
brain  
brawn  
prawn



formally:  $\langle (train, lb, brain), (brain, 4w, brawn), (brawn, lp, prawn) \rangle$

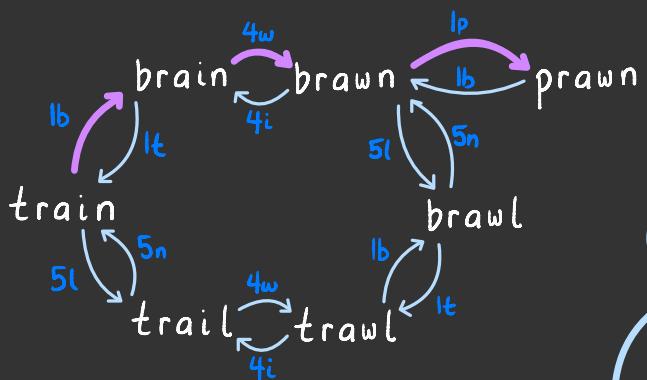
train  
trail  
trawl  
brawl  
brawn  
prawn



formally:  $\langle (train, 5l, trail), (trail, 4w, trawl), (trawl, lb, brawl), (brawl, 5n, brawn), (brawn, lp, prawn) \rangle$

# search paths

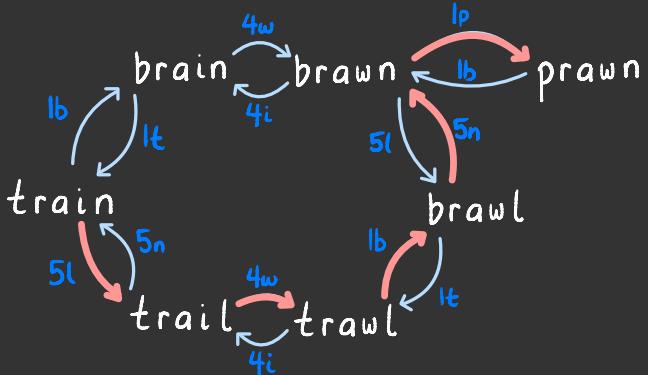
train  
brain  
brawn  
prawn



why do we include?  
the action.

formally:  $\langle (train, lb, brain), (brain, 4w, brawn), (brawn, lp, prawn) \rangle$

train  
trail  
trawl  
brawl  
brawn  
prawn

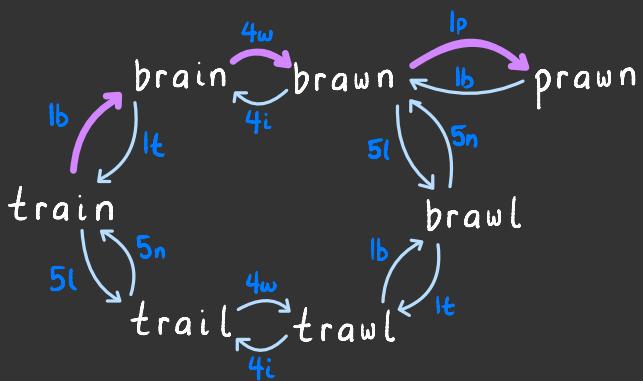


why do we include?  
the next state.

formally:  $\langle (train, 5l, trail), (trail, 4w, trawl), (trawl, lb, brawl), (brawl, 5n, brawn), (brawn, lp, prawn) \rangle$

# search paths

train  
brain  
brawn  
prawn



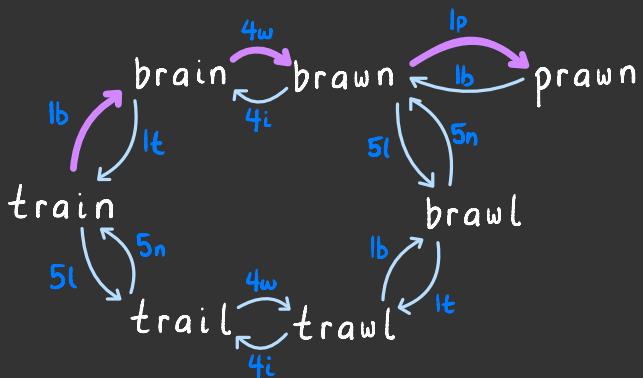
formally:  $\langle (train, lb, brain), (brain, 4w, brawn), (brawn, lp, prawn) \rangle$

states	$Q = \{train, brain, trail, \dots\}$
actions	$\Sigma = \{1a, \dots, 5z\}$
transitions	$\Delta = \{(train, lb, brain), \dots\}$
initial state	$q_0 = train$
final states	$F = \{prawn\}$
weight function	$\omega = \{\delta \mapsto 1 \mid \delta \in \Delta\}$

a **search path** is a sequence  $\langle \delta_0, \dots, \delta_k \rangle$  of **transitions** from  $\Delta$  such that there exist **states**  $q_0, \dots, q_{k+1} \in Q$  and **actions**  $\sigma_0, \dots, \sigma_k \in \Sigma$  such that  $\delta_i = (q_i, \sigma_i, q_{i+1}) \quad \forall i \in \{0, \dots, k\}$

# search paths

train  
brain  
brawn  
prawn



formally:  $\langle (train, lb, brain), (brain, 4w, brawn), (brawn, lp, prawn) \rangle$

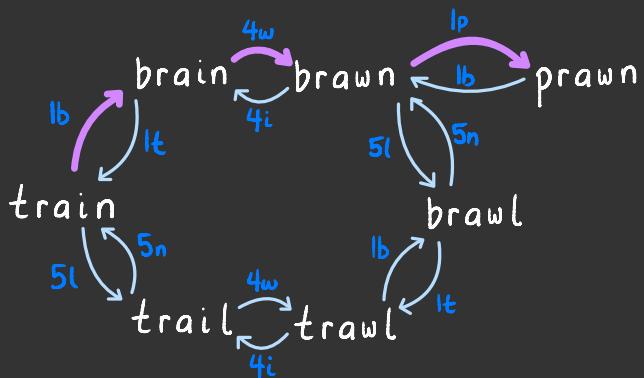
states	$Q = \{train, brain, trail, \dots\}$
actions	$\Sigma = \{1a, \dots, 5z\}$
transitions	$\Delta = \{(train, lb, brain), \dots\}$
initial state	$q_0 = train$
final states	$F = \{prawn\}$
weight function	$\omega = \{\delta \mapsto 1 \mid \delta \in \Delta\}$

a **search path** is a sequence  $\langle \delta_0, \dots, \delta_k \rangle$  of **transitions** from  $\Delta$  such that there exist **states**  $q_1, \dots, q_{k+1} \in Q$  and **actions**  $\sigma_0, \dots, \sigma_k \in \Sigma$  such that  $\delta_i = (q_i, \sigma_i, q_{i+1}) \quad \forall i \in \{0, \dots, k\}$

how does this differ from an arbitrary sequence of transitions?

# search paths

train  
brain  
brawn  
prawn



formally:  $\langle (train, lb, brain), (brain, 4w, brawn), (brawn, lp, prawn) \rangle$

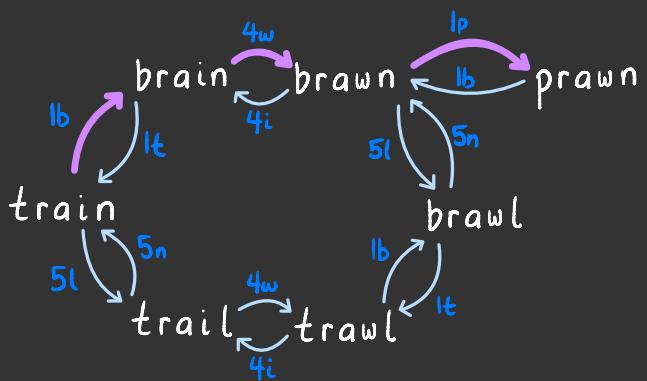
states	$Q = \{train, brain, trail, \dots\}$
actions	$\Sigma = \{la, \dots, 5z\}$
transitions	$\Delta = \{(train, lb, brain), \dots\}$
initial state	$q_0 = train$
final states	$F = \{prawn\}$
weight function	$\omega = \{\delta \mapsto 1 \mid \delta \in \Delta\}$

a search path is a sequence  $\langle \delta_0, \dots, \delta_k \rangle$  of transitions from  $\Delta$  such that there exist states  $q_0, \dots, q_{k+1} \in Q$  and actions  $\sigma_0, \dots, \sigma_k \in \Sigma$  such that  $\delta_i = (q_i, \sigma_i, q_{i+1}) \quad \forall i \in \{0, \dots, k\}$

why do these start at different indices?

# search paths

train  
brain  
brawn  
prawn



formally:  $\langle (train, lb, brain), (brain, 4w, brawn), (brawn, lp, prawn) \rangle$

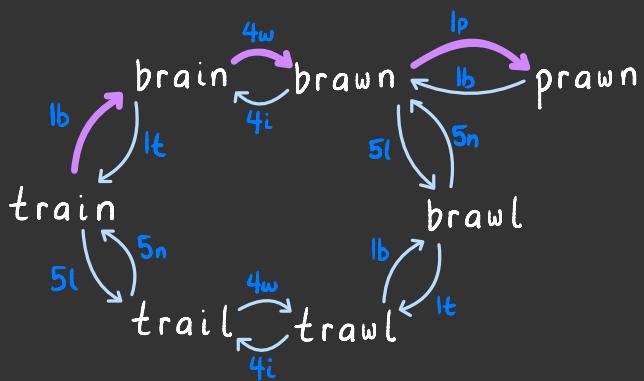
this must be the  
initial state

these must  
match

these must  
match

in a search path

train  
brain  
brawn  
prawn



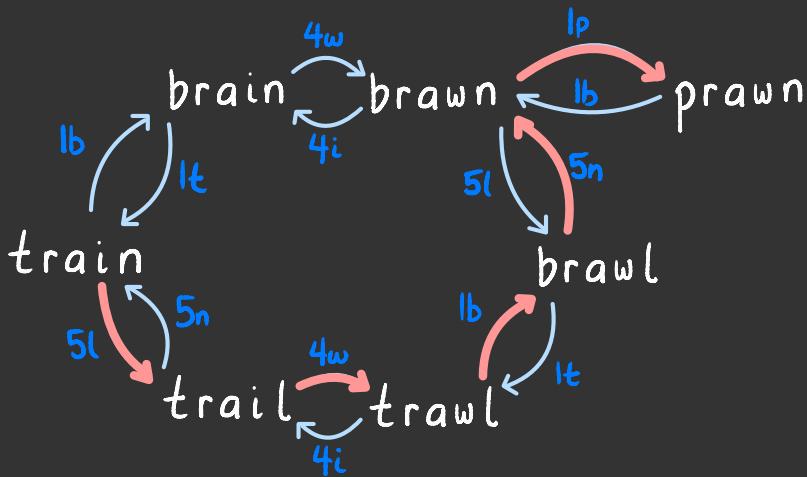
formally:  $\langle (train, lb, brain), (brain, 4w, brawn), (brawn, lp, prawn) \rangle$

states  $Q = \{train, brain, trail, \dots\}$   
actions  $\Sigma = \{la, \dots, 5z\}$   
transitions  $\Delta = \{(train, lb, brain), \dots\}$   
initial state  $q_0 = train$   
final states  $F = \{prawn\}$   
weight function  $\omega = \{\delta \mapsto 1 \mid \delta \in \Delta\}$

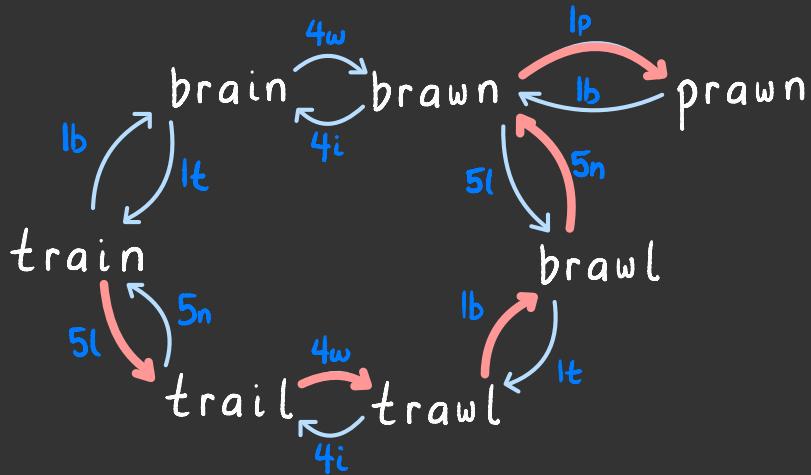
a search path is a sequence  $\langle \delta_0, \dots, \delta_k \rangle$  of transitions from  $\Delta$  such that there exist states  $q_0, \dots, q_{k+1} \in Q$  and actions  $\sigma_0, \dots, \sigma_k \in \Sigma$  such that  $\delta_i = (q_i, \sigma_i, q_{i+1}) \quad \forall i \in \{0, \dots, k\}$

the cost of a search path is the sum of the weights of the transitions

what is the cost of this search path?

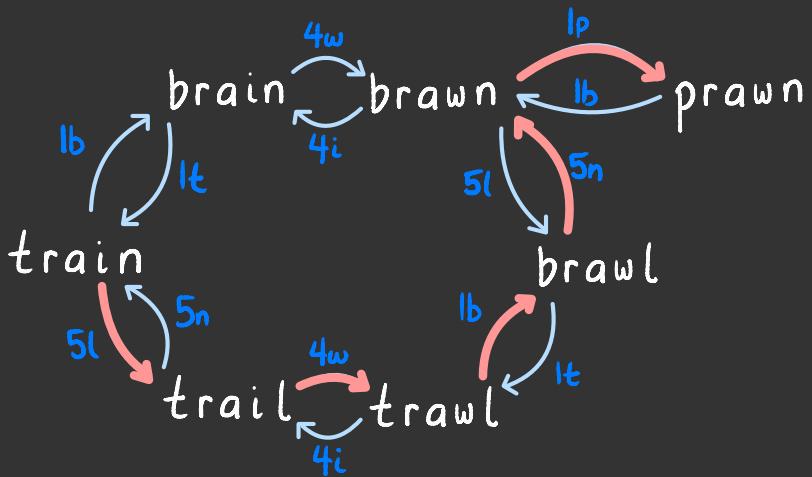


# what is the cost of this search path?



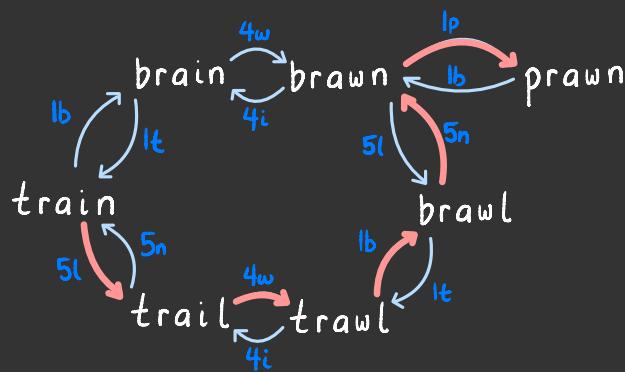
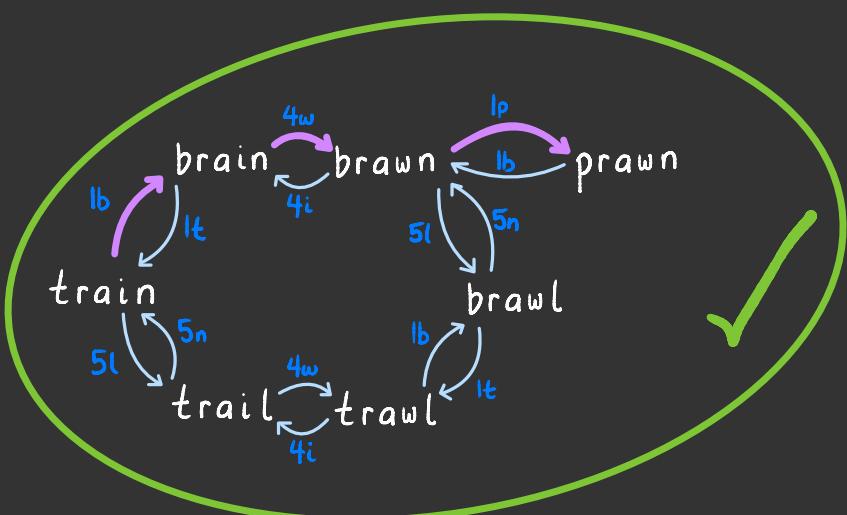
$$\begin{aligned} & w((train, 5l, trail)) \\ + & w((trail, 4w, trawl)) \\ + & w((trawl, lb, brawl)) \\ + & w((brawl, 5n, brawn)) \\ + & w((brawn, lp, prawn)) \end{aligned}$$

# what is the cost of this search path?

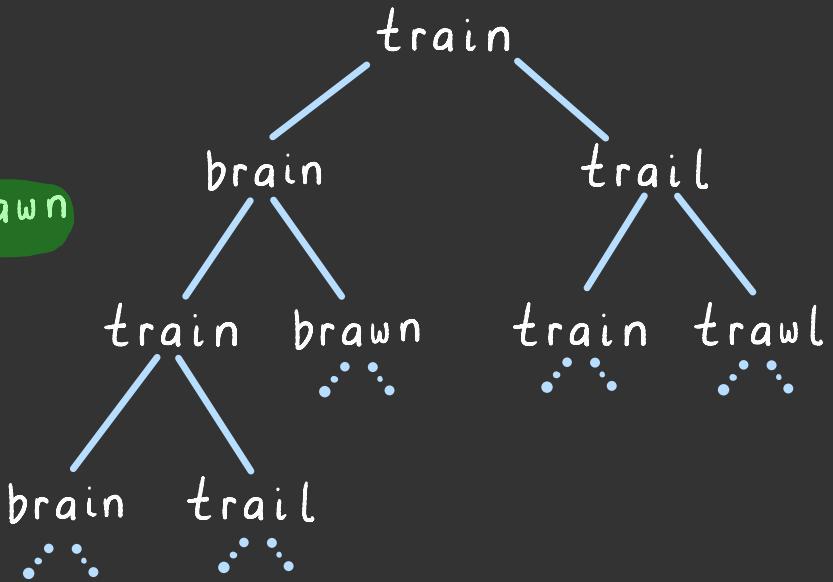
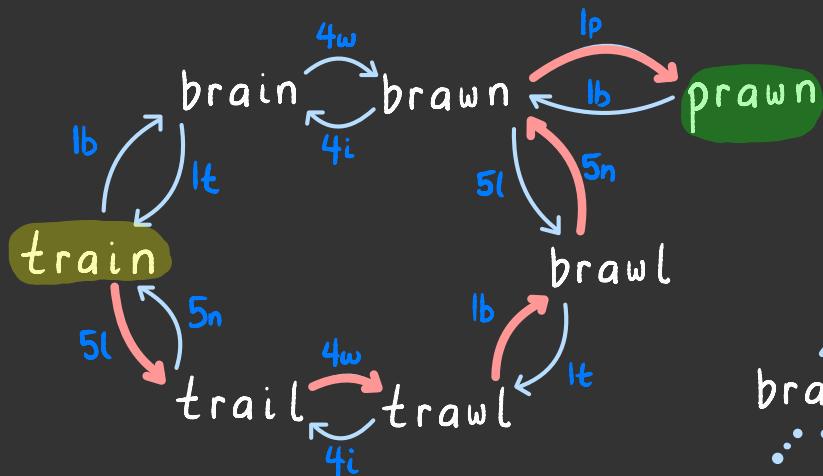


$$\begin{aligned} & w((train, 5l, trail)) & | \\ & + w((trail, 4w, trawl)) & + | \\ & + w((trawl, lb, brawl)) & + | \\ & + w((brawl, 5n, brawn)) & + | \\ & + w((brawn, lp, prawn)) & + | \\ & \hline & & 5 \end{aligned}$$

the goal of search:  
to find a **search path**  
of **minimal cost** whose  
last transition leads to  
a **final state**

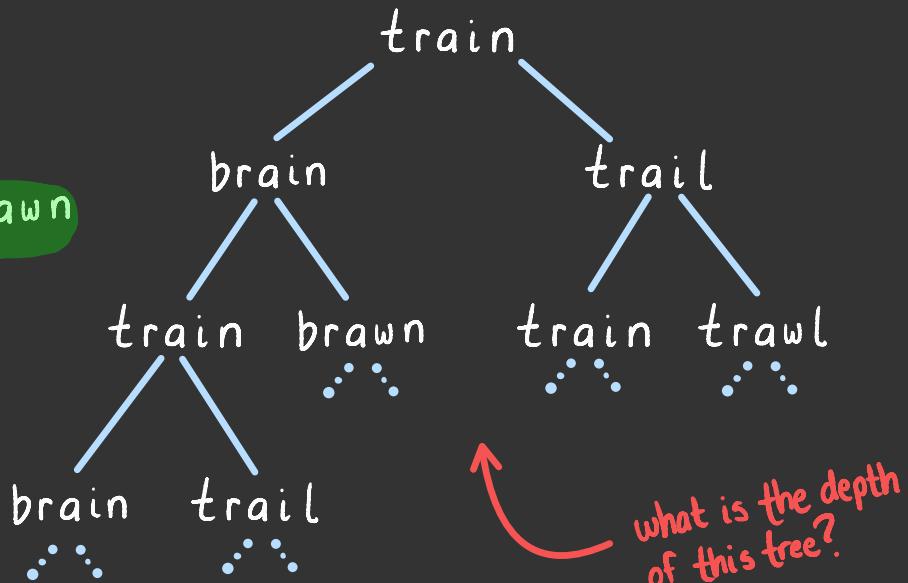
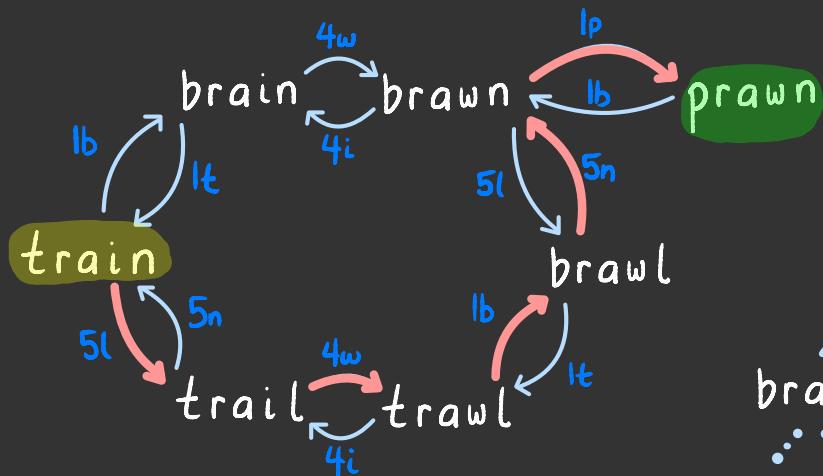


a state machine can be unraveled into  
a search tree



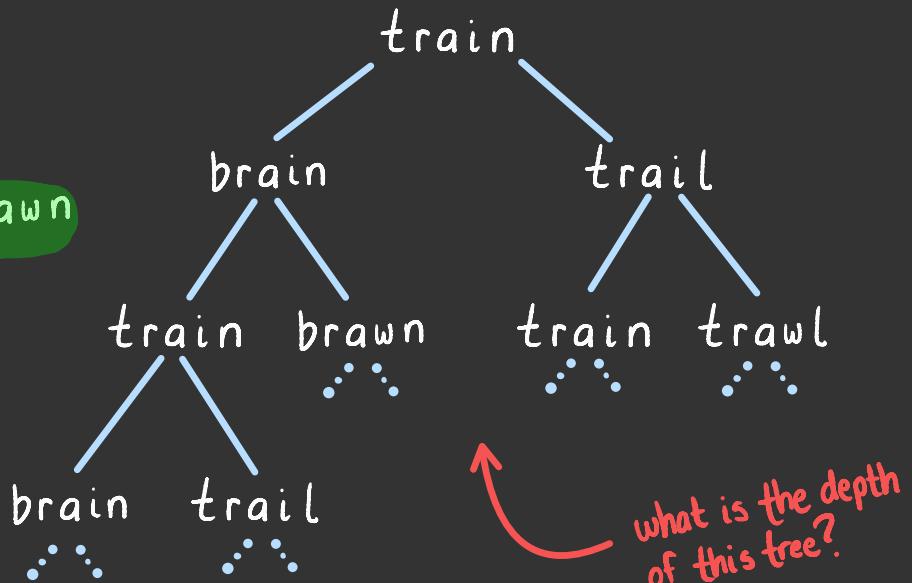
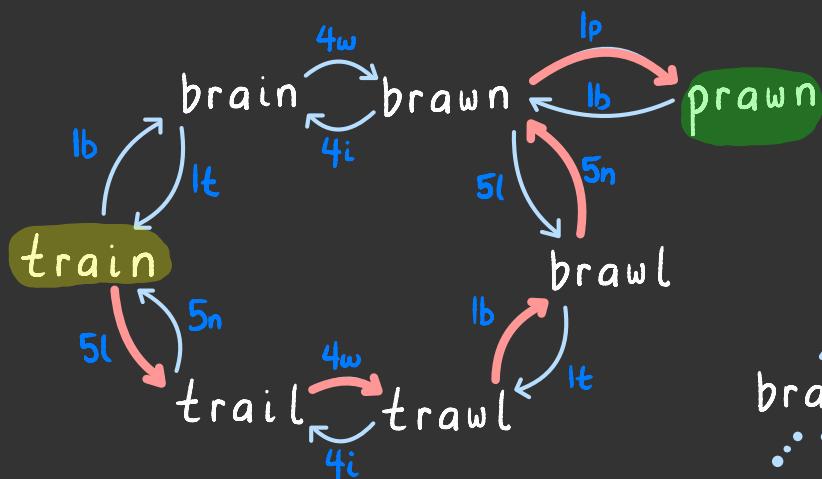
the search tree is a compact representation  
of all the search paths

a state machine can be unraveled into  
a search tree



your answer here

a state machine can be unraveled into  
a search tree



what is the depth  
of this tree?

infinite. in general, if a state machine has reachable cycles, then its search tree has infinite depth